

A Multi-Agent Governance Architecture for Human-Centric Digital Operations

White Paper Draft v0.1

Prepared by Filippo Forlani
June 3, 2026

Author and copyright notice

This document expresses the original conceptual framework proposed by Filippo Forlani. © 2026 Filippo Forlani. This work is licensed under the Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0).

This white paper is a technical and conceptual draft, not legal advice. Copyright normally protects the text, diagrams, and other original expressions of the model; recognition of the underlying idea is strengthened by dated publication, citation metadata, and archived releases.

Contents

1. Executive Summary	4
2. The Problem	4
2.1 Sources of risk.....	4
2.2 The authorization bottleneck	5
3. Current Models and Their Limits.....	5
3.1 Single-agent systems	5
3.2 Human approval systems.....	5
3.3 Copilots	6
3.4 Structural limitation	6
4. The ForIA ^{nl} Governance Model	6
4.1 Name and notation	6
4.2 Architectural principle	7
5. Roles.....	7
5.1 Ego: the impulse to act.....	7
5.2 Alter-Ego: the independent critic	7
5.3 Super-Ego: governance and bounded execution.....	7
5.4 Deus: the person in control	7
6. Workflow	7
6.1 Operational sequence	7
6.2 Risk classes	8
6.3 Minimum decision record	8
7. Applications.....	8
8. Security by Design	8
8.1 Recommended controls	8
8.2 Failure modes to test.....	9
9. Future Developments.....	9
Conclusions	9
Appendix A. Name and Notation	11
Appendix B. Formal Definition of the ForIA ^{nl} Governance Model.....	11
Appendix C. Governance Principles	11
Appendix D. Risk Escalation Framework.....	12
Appendix E. Example Operational Scenarios.....	12
E.1 Email reply with confidential content.....	12
E.2 Routine software update	12
E.3 Autonomous procurement	12
E.4 Calendar management	12
Appendix F. False Positive / False Negative Optimization	13
Appendix G. Future Research Directions	13
Appendix H. GitHub Publication Package	13
H.1 Files	13
H.2 License	14

H.3 Citation..... 14

Appendix I. Version History and Authorship Record 14

References..... 14

1. Executive Summary

The **ForIA^{AI} Governance Model** is a proposed multi-agent governance architecture designed to enable safe, human-centric digital operations. The model was conceived and developed from direct experience with AI agents and real-world user needs, while explicitly addressing the security challenges, operational risks, and inherent limitations associated with autonomous AI systems.

It addresses a central limitation of current AI assistants and autonomous agents: the same model often interprets a goal, proposes actions, evaluates its own plan, and prepares execution, while the human user is repeatedly asked to approve potentially sensitive operations.

The model separates action proposal, independent critique, governance synthesis, and final human authority into four roles: **Ego**, **Alter-Ego**, **Super-Ego**, and at the top of decisions, **Deus**.

Ego wants to act.

Alter-Ego cannot act but can criticize (Ego impulses).

Super-Ego cannot freely invent operations; it can only execute actions that originate from Ego and are evaluated by Alter-Ego.

Finally, **Deus** is the human owner, who confirms, rejects, or modifies operations when risk is too high or when the system faces ambiguity, conflict, or irreversible consequences.

The framework is intended for personal computing, email, calendar management, software development, enterprise workflows, autonomous procurement, backup operations, and other digital activities where AI systems may operate with increasing levels of delegated authority. It is designed to reduce both unsafe autonomy and excessive authorization fatigue.

Core thesis

Human control should not be reduced to repeated permission pop-ups. Human control should be embedded in a multi-agent governance process where proposals, critiques, decisions, logs, and escalations are structurally separated.

2. The Problem

Autonomous agents are powerful but dangerous. They can update software, read and answer emails, modify files, test software, manage calendars, perform backups, interact with remote services, and even make purchases on behalf of a person or organization. These capabilities are useful precisely because they can reduce repetitive work. They are also risky because each action may affect data integrity, security, reputation, legal position, financial exposure, or relationships with third parties.

For this reason, current tools frequently require the user to grant explicit permission before performing sensitive actions on the operating system, email accounts, calendars, code repositories, remote services, or external platforms. This repeated authorization pattern is understandable: it keeps the human user formally in control. However, it also reveals a structural weakness. The system does not always contain a sufficient internal governance mechanism before it reaches the human.

2.1 Sources of risk

- **Hallucination risk.** The AI may infer facts, constraints, file paths, people, dates, or intentions incorrectly and then act on those incorrect assumptions.
- **Operational damage.** Commands can delete, overwrite, move, expose, or corrupt data; system updates can break compatibility; automatic configuration changes can destabilize a computer or network.
- **Reputational damage.** An AI-generated email, calendar invitation, message, or document may be factually wrong, poorly phrased, too informal, legally imprudent, or simply inconsistent with the person it represents.
- **Financial exposure.** Autonomous procurement or paid online actions may create unintended expenses or contractual obligations.
- **Security exposure.** The AI may upload files, send sensitive information, execute untrusted code, or interact with malicious content.
- **Lack of personal sensitivity.** AI systems may be highly competent at synthesis and repetition, but they lack the lived judgement, social nuance, risk appetite, and contextual sensitivity of the person they assist.

A further issue is **personality alignment**. If an AI is expected to replace a person in repetitive or synthetic tasks, it should not merely be intelligent; it should also act in a way that is compatible with the person's preferences, habits, tone, caution level, professional role, and tolerance for risk. Even if such a personal sensitivity were technically approximated, it would still require governance: a system that behaves like the user must not be free to act without limits.

2.2 The authorization bottleneck

The common response to these risks is to ask for permission repeatedly. This creates an authorization bottleneck: the user is interrupted to approve each step, often without enough structured information to judge the risk efficiently. Over time, repeated prompts may cause approval fatigue. The user may approve too quickly, reject useful automation, or lose trust in the system.

The key problem is therefore not only whether the user approves or rejects an action. The key problem is how the system reaches the point where the user is asked.

3. Current Models and Their Limits

Current AI-assisted operating patterns can be simplified into three families: single-agent autonomy, human approval workflows, and copilots. These models are useful, but they often retain the same basic architecture: the AI proposes an action, the human approves or rejects, and the system executes.




Figure 1. Current authorization loop in many single-agent, copilot, or tool-calling systems.

3.1 Single-agent systems

A single-agent system receives a goal, reasons about it, selects tools, and may execute actions. Its strength is simplicity. Its weakness is concentration of functions: the same reasoning process may define the task, propose the action, evaluate the action, and justify the action. This creates a weak separation between intention, criticism, and execution.

3.2 Human approval systems

Human approval systems reduce risk by requesting permission before sensitive operations. However, approval is often reactive. The human is asked to intervene at the end of a chain of reasoning, rather than after a structured internal critique. Human approval remains necessary, but it should be reserved for meaningful decisions, exceptions, high-risk operations, and policy conflicts, not for every minor step.

3.3 Copilots

Copilots are designed to assist rather than fully replace the user. They are effective for drafting, coding, summarizing, searching, and proposing actions. Their limitation appears when they become increasingly operational: a copilot that can act on files, calendars, email, cloud services, operating systems, and purchases requires stronger governance than ordinary autocomplete or drafting assistance.

3.4 Structural limitation

The common structural limitation is that authorization is treated as an external user prompt rather than as the final step of an internal governance process. The human remains responsible, but the system does not always provide an independent, role-separated evaluation before escalation.

4. The ForIAⁿI Governance Model

The **ForIAⁿI Governance Model** proposes a flexible, agile framework based on multiple AI agents, each with a constrained role. The model is not merely a collection of agents. Its defining feature is role separation:

- the agent that proposes does not execute;
- the agent that critiques cannot act;
- the agent that executes cannot freely invent operations;
- and the human remains the final authority when risk, conflict, or uncertainty exceeds the system's delegated envelope.

4.1 Name and notation

The name ForIAⁿI combines the acronym AI with a stylized rendering of the author's surname, Forlani. Beyond its graphical identity, the notation conveys an architectural concept.

The superscript *n* denotes extensibility and scalability. While the reference configuration described in this paper is based on four governance roles—Ego, Alter-Ego, Super-Ego, and Deus—the framework is not restricted to this configuration. Additional specialized agents may be introduced to address specific domains such as cybersecurity, legal compliance, financial oversight, privacy protection, or technical validation.

Regardless of the number of agents involved, the architecture preserves its fundamental governance principles: separation of proposal and execution, independent critique, bounded and auditable execution, risk-based escalation, and preservation of human sovereignty.

The notation may also be interpreted as 'AI to the *n*th power', emphasizing the layered and scalable nature of the governance process. Throughout this document, ForIAⁿI Governance Model (FGM) is used as the formal designation of the proposed framework.

4.1.1 Psychological Inspiration and Terminology

The terms **Ego**, **Alter-Ego**, and **Super-Ego** are intentionally inspired by concepts commonly associated with psychology and psychoanalytic literature. Their use within the ForIAⁿI Governance Model is primarily conceptual and pedagogical.

The framework does not attempt to reproduce, simulate, validate, or extend any specific psychological theory. Instead, these terms are employed as intuitive metaphors that help describe the functional roles of the different governance layers.

- **Ego** represents the impulse to act and pursue objectives.
- **Alter-Ego** represents an independent critical perspective capable of questioning proposed actions.
- **Super-Ego** represents governance, synthesis, and controlled execution.
- **Deus** represents the human decision-maker and ultimate authority.

The terminology was chosen because it provides a familiar conceptual vocabulary that facilitates understanding of role separation and decision governance within a multi-agent AI architecture. The model should therefore be interpreted as an engineering and governance framework rather than as a psychological model.

4.2 Architectural principle

The model adopts a principle similar to separation of duties: a system is safer when proposal, critique, decision, execution, and final authority are not concentrated in a single component. The goal is not to eliminate human authorization. The goal is to make human authorization less frequent, more meaningful, better documented, and supported by a structured internal evaluation.

5. Roles

Role	Formal function	Capability	Hard limit
Ego	Autonomous Proposal Engine	Proposes operational actions and plans.	Cannot execute directly.
Alter-Ego	Independent Risk & Critique Engine	Performs independent critique and risk assessment.	Cannot execute or propose new operations as actions.
Super-Ego	Governance & Execution Engine	Decides and executes bounded operations.	Cannot execute actions not originating from Ego and not evaluated by Alter-Ego.
Deus	Human Authority & Override	Final approval, rejection, override, or modification.	Human authority; not automated.

5.1 Ego: the impulse to act

Ego is the operative impulse of the system. It identifies possible actions that may satisfy the owner's goals. Ego may be proactive, but its output must remain propositional. It can say: "I suggest updating this software," "I suggest replying to this email," or "I suggest backing up this folder." It cannot directly perform those actions.

5.2 Alter-Ego: the independent critic

Alter-Ego is a distinct AI model or agent whose purpose is critical evaluation. It should not share the same operational impulse as Ego. It asks what could go wrong, whether the proposal is reversible, whether information is missing, whether the action could harm the user, and whether escalation to the human is required.

5.3 Super-Ego: governance and bounded execution

Super-Ego synthesizes the two previous outputs. It is the only automated component allowed to execute operations, and only within the policy envelope. Its authority is bounded: no proposal from Ego means no execution; no critical evaluation from Alter-Ego means no execution; high risk or unresolved conflict means escalation to Deus.

5.4 Deus: the person in control

Deus is the human owner. The name deliberately emphasizes that final authority is not delegated to the machine. Deus is not asked to micromanage every minor operation. Deus is called when there is significant risk, ambiguity, irreversibility, reputational exposure, financial exposure, or policy conflict.

6. Workflow

Figure 2. ForIA^{AI} workflow: proposal, critique, synthesis, execution, and escalation.

6.1 Operational sequence

1. The human defines goals, preferences, constraints, and standing policies.
2. Ego generates one or more candidate actions and explains their expected benefit.
3. Alter-Ego independently evaluates each candidate action, including risk, reversibility, missing information, and worst-case outcomes.
4. Super-Ego compares Ego's proposal with Alter-Ego's critique and checks the operation against policy.
5. If the action is low-risk, reversible, policy-compliant, and adequately justified, Super-Ego executes it and records the evidence.

6. If the action is high-risk, irreversible, financially relevant, reputationally sensitive, legally sensitive, ambiguous, or disputed, Super-Ego does not execute and escalates to Deus.
7. Deus approves, rejects, modifies, or postpones the operation.
8. The system records the proposal, critique, decision, execution logs, human override, and lessons learned.

6.2 Risk classes

These definitions are preliminary and should be calibrated through real-world testing and domain-specific validation.

Risk class	Examples	Default decision	Human escalation
Low	Summaries, drafts, update checks, backup plans.	Prepare only; no destructive action.	Normally no.
Medium	Routine updates, file moves, low-sensitivity replies, calendar edits.	Execute if reversible, logged, and policy-compliant.	Sometimes.
High	Data deletion, system changes, sensitive emails, purchases.	Do not execute automatically.	Yes.
Critical	Irreversible, financial, legal, reputational, or credential exposure.	Block automatic execution.	Always.

6.3 Minimum decision record

- User goal or trigger.
- Ego proposal, including intended action and expected benefit.
- Alter-Ego critique, including risks, missing information, and worst-case scenario.
- Super-Ego decision and policy basis.
- Execution status, timestamp, tool used, affected files/accounts/services, and rollback path.
- Deus decision when escalation occurs.

7. Applications

- **PC management.** Software updates, system cleanup, driver checks, configuration review, backup scheduling, and health monitoring.
- **Email management.** Email triage, reply drafting, attachment review, priority classification, and safe sending under explicit policies.
- **Calendar management.** Meeting scheduling, conflict detection, reminders, travel buffers, agenda preparation, and rescheduling proposals.
- **Software development.** Code review, testing, dependency updates, build automation, issue triage, pull-request preparation, and release checks.
- **Enterprise workflows.** Document routing, reporting, approvals, task delegation, knowledge-base updates, CRM/ERP interactions, and audit trails.
- **Autonomous procurement.** Price comparison, vendor checks, purchase proposal, budget verification, risk review, and human authorization before commitment.
- **Security.** Credential hygiene reminders, suspicious email analysis, access review, log monitoring, and controlled incident response escalation.

8. Security by Design

The ForIA^{AI} model should be implemented with explicit safety constraints rather than merely relying on good prompts. The architecture is intended to make unsafe behavior less likely by design.

8.1 Recommended controls

- **Least privilege.** Each agent receives only the permissions needed for its role. Alter-Ego should not have execution rights. Ego should not have execution rights. Super-Ego receives constrained execution rights.
- **Tool allowlists.** Sensitive tools should be available only through Super-Ego and only under policy.

- **Reversibility scoring.** Operations should be classified according to whether they can be undone.
- **Evidence logging.** Every proposal, critique, decision, and operation should be recorded in a human-readable log.
- **Human escalation thresholds.** High-risk, irreversible, financial, legal, reputational, or ambiguous operations should escalate to Deus.
- **Model diversity.** Ego and Alter-Ego may use different models, prompts, or evaluation procedures to reduce correlated failure.
- **Rollback strategy.** Where possible, execution should include a pre-action snapshot, backup, version control commit, or undo plan.

8.2 Failure modes to test

- Ego proposes an action based on an incorrect assumption.
- Alter-Ego misses a hidden risk.
- Super-Ego executes outside policy.
- The user approves too quickly because of prompt fatigue.
- The system logs are incomplete or difficult to audit.
- A malicious email, file, or webpage manipulates the agent through prompt injection.
- A low-risk action becomes high-risk because of context not visible to the AI.

9. Future Developments

The framework can evolve in several directions. The exponent n in ForIAⁿI indicates that additional specialized agents may be added without changing the underlying governance logic.

- **Policy memory.** A controlled memory of the owner's preferences, recurring decisions, style, risk tolerance, and authorization rules.
- **Domain-specific critics.** Specialized Alter-Ego agents for cybersecurity, legal exposure, finance, professional reputation, privacy, and technical reliability.
- **Simulation before execution.** A sandbox in which Ego's proposed action is simulated before Super-Ego can execute it.
- **Audit dashboard.** A user interface showing proposed actions, critiques, approvals, logs, and escalation history.
- **Delegation levels.** Explicit levels of autonomy: observe only, propose only, prepare but do not execute, execute reversible actions, execute within budget, escalate critical actions.
- **Enterprise policy integration.** Integration with role-based access control, compliance policies, data-loss prevention, security operations, and procurement rules.
- **Formal verification.** For high-risk environments, policy rules could be formally encoded and checked before execution.
- Development of differentiated ethical, behavioral, and risk-assessment models for Ego, Alter-Ego, and Super-Ego, with the goal of minimizing unnecessary Deus escalations (false positives) while maintaining a near-zero false-negative rate, ensuring that actions requiring human authorization are never executed autonomously.

Conclusions

The growing ability of AI agents to act directly on digital systems requires a governance model that is stronger than repeated permission prompts and more flexible than full prohibition. The ForIAⁿI Governance Model proposes a clear structure: Ego proposes, Alter-Ego critiques, Super-Ego decides and executes within policy, and Deus remains the final authority.

The model is intentionally simple. Its value lies in separating roles that are often merged: desire to act, critical evaluation, operational execution, and human authority. This separation can reduce unsafe autonomy, reduce authorization fatigue, and provide clearer responsibility when AI systems manage computers, emails, calendars, code, enterprise workflows, and external transactions.

The next practical step is to publish this draft with a clear authorship statement, diagrams, versioned releases, citation metadata, and a license that matches the author's intention: open recognition, controlled reuse, or full reservation of rights.

June 2026

Filippo Forlani

Appendix A. Name and Notation

The name ForIAⁿI combines the acronym AI with a stylized rendering of the author's surname, Forlani. The notation is intended both as a graphical identity and as a conceptual marker for a scalable governance architecture.

The superscript n expresses extensibility: the architecture may include more than four agents in specialized implementations, while preserving the same governing logic. A cybersecurity implementation, for example, may add a dedicated security critic; a financial implementation may add a budget-control critic; an enterprise implementation may add compliance and data-protection agents.

The notation also evokes the idea of 'AI to the nth power', meaning that the intelligence of the system is not concentrated in a single agent but distributed across multiple constrained roles. The preferred formal designation in this paper is ForIAⁿI Governance Model (FGM).

Appendix B. Formal Definition of the ForIAⁿI Governance Model

The ForIAⁿI Governance Model can be formally described as a role-separated multi-agent governance architecture for digital operations. Each component has a defined capability envelope and a defined prohibition envelope.

Role	Primary function	Can propose	Can execute	Hard constraint
Ego	Generate candidate actions and operational plans.	Yes	No	May not execute actions directly.
Alter-Ego	Perform independent critique, risk assessment, and impact analysis.	No, except as critique.	No	May not operate on external systems.
Super-Ego	Synthesize proposal and critique; execute bounded operations.	No autonomous origination.	Yes, within policy.	May only execute actions originating from Ego and evaluated by Alter-Ego.
Deus	Human authority, final approval, rejection, override, or modification.	Yes	Yes, as human owner.	Not automated; preserves human sovereignty.

The architecture is therefore defined less by the number of models used than by the separation of permissions. The critical requirement is that no single automated component should simultaneously originate, approve, and execute sensitive actions.

Appendix C. Governance Principles

- **Separation of duties.** Proposal, critique, decision, and execution should not be concentrated in the same agent.
- **Independent critique.** The critical review must be structurally independent from the impulse to act.
- **Bounded execution.** Execution must occur only within explicitly authorized tools, policies, and risk limits.
- **Human sovereignty.** The human owner remains the final authority for high-risk, ambiguous, irreversible, financial, legal, or reputational actions.
- **Explainability.** Every executed or rejected action should be supported by a concise rationale.
- **Auditability.** The system should log proposals, critiques, decisions, executions, escalations, and overrides.
- **Least privilege.** Each agent should receive only the permissions required by its role.
- **Risk-based escalation.** Human intervention should be required by risk and ambiguity, not by an undifferentiated prompt mechanism.

Appendix D. Risk Escalation Framework

The following matrix is a preliminary operational framework. It should be adapted to the owner, organization, legal environment, and domain of use.

Risk class	Typical action	Reversibility	Default route	Deus escalation
Low	Summaries, draft preparation, update checks, file listing.	High	Super-Ego may approve preparation or non-destructive execution.	Normally no.
Medium	Calendar edits, routine file organization, reversible software updates, low-sensitivity replies.	Medium to high	Super-Ego may execute if reversible, logged, and policy-compliant.	Conditional.
High	External emails, purchases, system configuration changes, deletion, sensitive documents.	Low to medium	Super-Ego should normally defer or request Deus approval.	Usually yes.
Critical	Irreversible, legal, financial, credential, security, or reputational exposure.	Low or none	Automatic execution blocked.	Always.

Risk classification should consider impact, probability of error, reversibility, exposure to third parties, sensitivity of data, financial value, legal relevance, and the user's personal or professional risk tolerance.

Appendix E. Example Operational Scenarios

Some examples of workflow.

E.1 Email reply with confidential content

- Ego proposes a reply to a supplier or client.
- Alter-Ego detects that the message may disclose confidential information or an unverified commitment.
- Super-Ego classifies the action as high risk because it affects reputation and external communication.
- Deus is asked to approve, modify, or reject the outgoing email.

E.2 Routine software update

- Ego proposes a routine update of a known software package.
- Alter-Ego checks compatibility, rollback options, and whether the update is from a trusted source.
- Super-Ego executes only if the update is reversible, logged, and consistent with the user's maintenance policy.
- Deus is not involved unless the update affects critical software or lacks rollback evidence.

E.3 Autonomous procurement

- Ego proposes purchasing replacement hardware or software.
- Alter-Ego evaluates price, vendor reliability, budget threshold, contractual implications, and alternatives.
- Super-Ego blocks automatic execution if the purchase exceeds delegated authority.
- Deus approves or rejects the purchase before any financial commitment is made.

E.4 Calendar management

- Ego proposes a calendar modification or meeting rescheduling.
- Alter-Ego evaluates conflicts, travel buffers, participant sensitivity, and the risk of appearing unprofessional.
- Super-Ego may execute low-risk internal scheduling changes if policies permit.
- Deus is involved for high-importance meetings, external stakeholders, or ambiguous social context.

Appendix F. False Positive / False Negative Optimization

A central design objective of the ForIAⁿI Governance Model is to optimize the trade-off between autonomy and safety. The model should reduce unnecessary human interruptions while maintaining rigorous control over actions that could harm the user, the system, or third parties.

Term	Definition in FGM	Consequence
False positive	The system escalates to Deus even though the action could have been safely handled within policy.	Unnecessary human interruption, reduced efficiency, approval fatigue.
False negative	The system executes autonomously an action that should have required Deus authorization.	Potential security, reputational, legal, financial, or operational damage.

The preferred safety strategy is asymmetric: false positives should be reduced, but not at the cost of increasing false negatives. In safety-critical or reputation-sensitive domains, a false negative is substantially more serious than a false positive.

The practical objective is therefore to minimize unnecessary Deus escalations while maintaining a near-zero false-negative rate, ensuring that actions requiring human authorization are never executed autonomously.

Appendix G. Future Research Directions

- **Differentiated ethical models for Ego, Alter-Ego, and Super-Ego**, so that each role applies a distinct reasoning style consistent with its function.
- **Quantitative risk scoring** based on impact, probability, reversibility, data sensitivity, financial exposure, and external visibility.
- **Policy learning mechanisms** that adapt to the user’s recurring decisions without removing final human authority.
- **Model diversity strategies** to reduce correlated failures between Ego and Alter-Ego.
- **Formal verification** of high-risk policies before execution.
- Sandbox simulation for operating-system changes, software updates, code execution, and external communications.
- Enterprise integration with role-based access control, compliance frameworks, data-loss prevention tools, and procurement workflows.

Appendix H. GitHub Publication Package

Repository name is:

`forlani-governance-model`

H.1 Files

File
README.md
white-paper/ForIA-n-I-Governance-Model.pdf
white-paper/ForIA-n-I-Governance-Model.docx
diagrams/current-authorization-loop.png
diagrams/forlani-governance-workflow.png
CITATION.cff
LICENSE.md
CHANGELOG.md

H.2 License

The ForIA^{AI} Governance Model is released under the **Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0)**.

This license permits:

- sharing and redistribution of the work in any medium or format;
- adaptation, modification, and extension of the work;
- citation and reuse in research, educational, and non-commercial projects;

provided that:

- appropriate credit is given to the original author;
- a link to the license is provided;
- any modifications are clearly indicated;
- the work is not used for commercial purposes without explicit authorization from the author.

The objective of this licensing strategy is to encourage discussion, academic citation, experimentation, and community-driven development of the framework while preserving the author's attribution rights and preventing unauthorized commercial exploitation.

License reference:

<https://creativecommons.org/licenses/by-nc/4.0/>

H.3 Citation

Forlani, F. (2026). The ForIA^{AI} Governance Model: A Multi-Agent Governance Architecture for Human-Centric Digital Operations. White paper draft v0.1.

Appendix I. Version History and Authorship Record

Version	Date	Author	Notes
v0.1	June 3, 2026	Filippo Forlani	Initial white paper draft and core workflow diagrams.

Concept and architecture: Filippo Forlani. This record is intended to support dated attribution, version control, and future citation of the ForIA^{AI} Governance Model.

References

- GitHub Docs. Creating a new repository. <https://docs.github.com/en/repositories/creating-and-managing-repositories/creating-a-new-repository>
- GitHub Docs. About READMEs. <https://docs.github.com/en/repositories/managing-your-repositorys-settings-and-features/customizing-your-repository/about-readmes>
- GitHub Docs. Licensing a repository. <https://docs.github.com/repositories/managing-your-repositorys-settings-and-features/customizing-your-repository/licensing-a-repository/>
- GitHub Docs. About CITATION files. <https://docs.github.com/en/repositories/managing-your-repositorys-settings-and-features/customizing-your-repository/about-citation-files>
- GitHub Docs. Managing releases in a repository. <https://docs.github.com/en/repositories/releasing-projects-on-github/managing-releases-in-a-repository>
- Zenodo Help. GitHub and Software. <https://help.zenodo.org/docs/github/>
- Creative Commons. Attribution-NonCommercial-NoDerivatives 4.0 International. <https://creativecommons.org/licenses/by-nc-nd/4.0/>